

Tournaments not inducible by five voters

Leonid Chindelevitch¹ and Ararat Harutyunyan²

¹ MRC Centre for Global Infectious Disease Analysis, School of Public Health, Imperial College London, London, United Kingdom. Email: 1chindel@ic.ac.uk. ORCID: 0000-0002-6619-6013.

² University of Paris-Dauphine, PSL University, CNRS UMR7243, LAMSADE, Paris, France. Email: ararat.harutyunyan@lamsade.dauphine.fr.

Abstract

A tournament T is k -inducible if there are k linear orders on its vertex set such that, for every arc $i \rightarrow j$ of T , a majority of the orders rank i above j . For odd k , let $N(k)$ be the least order at which some tournament is not k -inducible. Only $N(3) = 8$ is known exactly; for $N(5)$ the best bounds were $12 \leq N(5) \leq 38$, from our previous paper [2], which also gave the first explicit example of moderate order, the Paley tournament P_{43} .

Results. A bespoke search algorithm improves both ends: $13 \leq N(5) \leq 23$. The upper bound comes from proving that P_{23} is not 5-inducible, the case Bachmeier et al. [1] reported they could not decide, their SAT solver not having terminated within a cumulative six weeks; ours takes 22 hours on one laptop. The lower bound comes from an analysis at order 12. We also show that P_{31} is not 5-inducible, while P_{19} is 5-inducible but not with unit margin, that is, not by a profile in which every arc is carried by exactly three voters against two. Both P_{19} and P_{23} are arc-critical for their respective properties, whereas P_{31} and P_{43} are not vertex-critical: deleting a vertex leaves a tournament that is still not 5-inducible.

Method. The search places one vertex at a time, always choosing the vertex with the fewest options left, and propagates the consequences. Together with the automorphisms of the tournament, this decides on a single laptop instances that neither integer programming nor a general-purpose SAT solver can settle. The refutations for P_{19} and P_{23} are certified as well: the search is split into independent subproblems, a SAT solver emits a machine-checkable proof for each, and a separate program rechecks every proof. All results, subject to two human-checked lemmas, are reproducible from <https://github.com/Leonardini/TournamentsBeyond5Voters>.

1. Introduction

We follow the notation of our previous paper [2] throughout, which we briefly restate here. Let T be a *tournament* on the vertex set V , $|V| = n$: a digraph in which every unordered pair $\{i, j\}$ carries exactly one of the arcs $i \rightarrow j$ or $j \rightarrow i$. A k -*profile* consists of k linear orders $\pi_1, \dots, \pi_k$ on V . The *majority tournament* of such a profile puts an arc from i to j whenever $|\{v : i \prec_{\pi_v} j\}| > k/2$. For odd k this is well-defined on every pair, and the result is a tournament. We say T is k -*inducible* if some k -profile has T as its majority tournament, and we call such a profile a *witness* for T . The *converse* of T is the tournament obtained from T by reversing the direction of every arc.

We call the number of voters who rank i above j the *support* of the arc $e : i \rightarrow j$, and denote it $c(e)$. If $c(e) > k/2$, the *margin* of e is $2c(e) - k$. In a witness every arc's support exceeds $k/2$, so for $k = 5$ it is 3, 4 or 5. A witness has *margin at most* M if every arc has margin at most M . For $k = 5$ the possible margins are 1, which we call *unit margin*, 3 (every support in $\{3, 4\}$), and 5 (no restriction, i.e. plain majority). Note that a unit-margin witness is in particular a majority witness. In [2] we conjectured that the margin-based hierarchy is strict, which for $k = 5$ asserts two separate things: that some tournament is 5-inducible with margin at most 3 but not with unit margin, and that some tournament is 5-inducible but not with margin at most 3. We settle the first here since P_{19} is such a tournament (Section 3.3), while the second remains open (Section 4.2).

One further quantity recurs below. Informally, the *predictability* $\alpha^*(T)$ is the largest fraction $\alpha > 1/2$ for which there is some profile, on any number t of voters, such that every arc of T is supported by at least αt of them; it is the value of a linear program over the linear orders of V . Since every arc of a k -voter witness is supported by more than $k/2$ of the voters, k -inducibility forces $\alpha^*(T) \geq (k+1)/2k$, which is $\frac{2}{3}$ at $k=3$ and $\frac{3}{5}$ at $k=5$. The implication is only one-directional: α^* can meet the threshold and the tournament still fail to be inducible. Shepardson and Tovey [7] introduced the notion; see Section 1.4 in [2] for the computations and the counterexamples at both $k=3$ and $k=5$.

Counting arguments show that almost all tournaments fail to be k -inducible for any fixed k , yet the *explicit* examples they yield are astronomically large. This paper closes much of the gap for $k=5$. Let $N(k)$ be the least order of a tournament that is not k -inducible. The answer is only known exactly for $k=3$. For $k=5$ the state of the art before this work was

$$12 \leq N(5) \leq 38,$$

both ends of which are due to our previous paper [2]: the lower bound came from an exhaustive census of all $\approx 9.0 \times 10^8$ tournaments on 11 vertices, all of which admit a unit-margin five-voter profile, improving on Bachmeier et al.'s $N(5) \geq 11$ [1].

The upper bounds in the literature are of two different kinds and they are not always comparable. Some are non-constructive, and obtained by counting: one compares the number of tournaments of a given order with the number that k linear orders can produce, and when the first exceeds the second, some tournament of that order must be missed. Bachmeier et al. [1] argued this way to reach 41, and [2] sharpened the comparison by restricting it to regular and near-regular tournaments, reaching 38. Other bounds are constructive, and exhibit an explicit tournament. The record for $k=5$ was P_z on $z=603,979,799$ vertices [1], from a dominating-set theorem of Alon et al. [4] together with a construction of Graham and Spencer [5], which we improved to P_{43} [2], the first example of moderate order. Our bespoke algorithmic decision procedure shows that P_{23} is not 5-inducible, which improves both the best non-constructive (counting) as well as the best constructive (explicit) upper bounds to $N(5) \leq 23$.

1.1 Special properties of Paley tournaments

The Paley tournament P_q , for a prime $q \equiv 3 \pmod{4}$, has vertex set $\mathbb{F}_q$ and an arc from i to j whenever $j-i$ is a nonzero square; as everywhere in this paper, vertices are numbered from 1, vertex $i+1$ being the field element i . It is *doubly regular*: every vertex has out-degree $(q-1)/2$ and every ordered pair of vertices has exactly $(q-3)/4$ common out-neighbours. Its automorphism group, written $\text{Aut}(P_q)$ and formally defined in Section 2, has order $q(q-1)/2$. This makes P_q *arc-transitive*, meaning that for every pair of arcs e, e' of P_q there is an automorphism of P_q that takes e into e' . It is also *self-converse*, meaning that it is isomorphic to its own converse; the relevant mapping is simply multiplication by -1 , which is easily checked to not be an automorphism when $q \equiv 3 \pmod{4}$.

1.2 Why Paley(23) is hard to classify and Paley(43) was easy

Paley tournaments are natural candidates for non-inducibility because they are far from being transitive, and because their symmetry makes an otherwise hopeless search merely very hard. Bachmeier et al. [1] call them *quadratic residue tournaments* and write Q_q ; their orientation is the converse of ours. Since reversing every voter's order reverses every arc of the majority tournament and changes no support, a tournament and its converse have the same *majority dimension*, the least k for which a tournament is k -inducible, and are inducible at exactly the same margins. We use the name Paley and the notation P_q throughout.

The SAT study in [1] establishes

$$\dim(Q_3) = \dim(Q_7) = 3, \quad \dim(Q_{11}) = \dim(Q_{19}) = 5,$$

where $\dim$ denotes majority dimension, and then stops exactly where the problem becomes hard:

“Unfortunately, we were not able to check whether the majority dimension of Q_{23} is equal to 5 or larger as the SAT solver did not terminate within a total of six weeks.”

By contrast, our previous proof that P_{43} is not 5-inducible [2] never searched the profile space in general. It rests on a quantity we call the *slack*,

$$\text{slack}_5(T) = 5 \text{MAS}(T) - 3C, \quad C = \binom{n}{2},$$

where MAS is the maximum number of arcs of T that a single linear order can agree with. Summing the majority condition over all arcs bounds how far the five voters can *collectively* fall short of that maximum. A small slack restricts the voters to be close to the best orders, and at $q = 43$, where $\text{slack}_5 = 6$, that restriction allowed us to fully enumerate about 1.8 million cases.

While the same argument is valid at every q , the slack is *not* a monotone function of q (table from [2]):

q	7	11	19	23	31	43
slack_5	7	10	22	46	30	6

Since P_{43} attains a locally minimal slack, while P_{23} attains a local maximum, the method of [2] is unusable at $q = 23$. Note that for any q ,

$$\text{slack}_5(P_q) = 5C(\alpha^*(P_q) - \frac{3}{5}),$$

where α^* is the predictability; equality holds because P_q is arc-transitive [2]. In particular, the necessary condition $\alpha^* \geq 3/5$ is far from binding, so no argument resting on the predictability can certify infeasibility.

Predictability separates $k = 5$ from $k = 3$. At $k = 3$ predictability decides exactly at $N(3)$: the 96 non-3-inducible tournaments on eight vertices have $\alpha^* = \frac{13}{20}$, which is *below* $\frac{2}{3}$, so the relaxation detects them, while all other tournaments on eight vertices or fewer with at least one cycle have $\alpha^* = 2/3$ [7], and $N(3) = 8$ follows. At $k = 5$ it does not: the table above has $\alpha^* > \frac{3}{5}$ at every q , yet P_{23} , P_{31} and P_{43} are all non-inducible, with P_{23} being the farthest *above* the threshold among the three. Though some tournament with order between 13 and 22 may have predictability below the $\frac{3}{5}$ threshold, we conjecture in Section 4 that this is not the case.

1.3 Contributions

1. **P_{23} is not 5-inducible** (Section 3.1). This settles the instance on which the SAT solver of [1] did not terminate within a cumulative six weeks, and improves the best upper bound on $N(5)$ from 38 to 23. We also decide P_{31} , which does not improve the bound but demonstrates the method’s reach: the larger tournament is settled in less time than P_{23} .
2. **P_{19} is 5-inducible, but not 5-inducible with unit margin** (Section 3.3). This is an example suggesting that $k = 5$ may qualitatively differ from $k = 3$. For $k = 3$, the first k -inducible tournaments that are non- k -inducible with unit margin appear after $N(k)$, order 9 vs order 8 [2]. If P_{19} and P_{23} are the minimal examples in their respective classes, which we provide evidence for in this paper, the reverse is true for $k = 5$.
3. **P_{19} and P_{23} are arc-critical** (Sections 3.2 and 3.5). Reversing an arc in P_{19} restores 5-inducibility at unit margin, and reversing an arc in P_{23} restores unrestricted 5-inducibility.
4. **P_{31} and P_{43} are not vertex-critical** (Section 3.4). Deleting a vertex from P_{31} or P_{43} leaves a tournament that is still not 5-inducible.
5. **Two formally certified negative results** (Appendix B). P_{19} is not 5-inducible at unit margin and P_{23} is not 5-inducible. The search is split into independent subproblems (*cubes*); each is refuted by the SAT solver CaDiCaL, which emits its refutation as a machine-checkable *LRAT* proof; and every one of those proofs is then rechecked by a separate program that shares no code with the solver.

6. **Every tournament on 12 vertices is 5-inducible** (Appendix C). The analysis settles all 2,048 one-vertex extensions of each of the $\approx 9.0 \times 10^8$ tournaments on 11 vertices, which is complete because every 12-vertex tournament is such an extension of each of its twelve one-vertex deletions.

1.4 Related work

Two papers stand immediately behind this one, in different ways. Our previous paper [2] is the prior art for the *bound*: it produced the first explicit non-5-inducible tournament of modest size, P_{43} , and its two bounds are what we measure against: $N(5) \leq 38$ by counting and $N(5) \leq 43$ explicitly. Bachmeier et al. [1] is the prior art for the *problem and the method*: it brought SAT to bear on k -majority digraphs, characterised the 3-majority case building on Dushnik and Miller [8], and established hardness results for voting with a constant number of voters.

The value of $N(3)$ is settled. Shepardson and Tovey [7] introduced the predictability $\alpha^*(T)$ defined above. A 3-voter profile is exactly a $\frac{2}{3}$ -supermajority representation, so $\alpha^*(T) < \frac{2}{3}$ rules out 3-inducibility; their 8-vertex tournament has $\alpha^* = \frac{13}{20}$, which gives $N(3) \leq 8$. They also prove that every tournament on at most 7 vertices has $\alpha^* \in \{\frac{2}{3}, 1\}$, and Eggermont, Hurkens and Woeginger [6] then show by direct computation using an ILP formulation that every tournament on 7 vertices is 3-inducible, and every tournament on 8 or 9 vertices is 5-inducible, yielding $N(3) = 8$ and $N(5) \geq 10$.

Our certified computation (Appendix B.1) is standard. It splits the problem into independent subproblems, using the *cube-and-conquer* paradigm [9]: a SAT instance is split by fixing a few variables at a time, giving many partial assignments called *cubes*, and each cube is handed to the solver as a problem of its own. We adopt the standard of evidence set by the Boolean Pythagorean triples proof [3, 11], in which every subproblem emits a machine-checkable refutation that an independent checker validates. Our refutations are in the LRAT format [10] (Appendix B.1).

2. Preliminaries

Throughout, T is a tournament on $V = \{1, \dots, n\}$ and $k = 5$ unless stated otherwise. We write $N^+(v)$ and $N^-(v)$ for the out- and in-neighbourhoods of v . A tournament is *regular* if $|N^+(v)| = |N^-(v)|$ for each $v \in V$; this can only happen for n odd. Following [2], D_n , R_n and S_n denote the numbers of isomorphism classes of, respectively, all tournaments, regular tournaments and self-converse tournaments on n vertices. The text gives them to two significant figures and Appendix D tabulates their exact values. For any $S \subseteq V$, the sub-tournament of T induced by S is the tournament obtained by restricting T to the vertices in S and the arcs between those vertices.

We call a tournament T *vertex-critical* for k -inducibility at margin M if T is not k -inducible at margin M , but $T - v$ is k -inducible at margin M for every vertex v . Similarly, we call T *arc-critical* for such a property if T has it, but any T' obtained from T by reversing an arc does not.

2.1 Base states

Every search starts by fixing how the voters rank a handful of vertices and then looks for ways to add the remaining ones. Choose a *base* $B \subseteq V$ of b vertices (we use $b = 5$ or $b = 6$), and write $T|_B$ for the sub-tournament induced on B . When B is a base, a *base state* of B (we omit B when the context makes it clear) is a k -voter profile over B that is consistent with $T|_B$, i.e. k orderings of B whose majorities reproduce $T|_B$ at the margin being asked for. Base states differing only by a relabelling of the voters are counted once (see Lemma 2.3 below).

Each base state of B defines a self-contained subproblem: the rankings of B are given, and the question is whether the remaining $n - |B|$ vertices can be inserted so that every arc of T has the right support, so

the search splits into independent components across base states. How many there are depends only on $|B|$, on $T|_B$ and on the margin regime, but *not* on n ; thus, with the five-vertex bases we use on the Paley tournaments, which induce regular sub-tournaments, there are 8,031 of them in both P_{23} and P_{47} .

2.2 Symmetry reductions

An *automorphism* of a tournament T is a permutation of its vertices carrying every arc to an arc; these form a group $\text{Aut}(T)$ under composition, and T is *rigid* when $\text{Aut}(T)$ is trivial. For a permutation group Γ on a set X , the *orbit* of x is $\Gamma x = \{\gamma(x) : \gamma \in \Gamma\}$, and the orbits partition X . This group is *transitive* on X when this partition is trivial (equivalently, for any $x, x' \in X$ there is a $g \in \Gamma$ with $g(x) = x'$). We use orbits of $\text{Aut}(T)$ on the vertices, on the arcs, and on the *ordered pairs* of distinct vertices, where γ sends (u, v) to $(\gamma(u), \gamma(v))$. An ordered pair is either an arc or the reverse of one, so that third action has at least two orbits.

Two symmetries act on witnesses and both are exploited by every method here: the $k!$ relabellings of the voters, and $\text{Aut}(T)$ acting on the vertices. Each yields a reduction without loss of generality. As a witness can be transformed by an automorphism, we may in some cases decide in advance what some voter's top two vertices are.

Lemma 2.1 (orbit anchoring). *Let $\text{Aut}(T)$ act on the ordered pairs of distinct vertices, and choose one representative pair from each orbit, say $(a_1, b_1), \dots, (a_m, b_m)$. If T is k -inducible, then T has a witness in which, for at least one index i , some voter ranks a_i first and b_i second.*

Proof. Let $\pi_1, \dots, \pi_k$ be a witness and let (u, w) be the top two vertices of π_1 . The pair (u, w) lies in one of the orbits; let (a_i, b_i) be that orbit's representative, so there is $g \in \text{Aut}(T)$ with $g(u) = a_i$ and $g(w) = b_i$. Applying g to every voter gives a profile whose majority tournament is $g(T) = T$, so it is again a witness, and its first voter ranks a_i first and b_i second. ■

Thus we may search only for witnesses whose top pair is one of the m chosen pairs. Note that the lemma only guarantees that *some* i works, not any particular one. For a Paley tournament $m = 2$, since an ordered pair is either an arc or a non-arc and $\text{Aut}(P_q)$ is transitive on each set. The two representatives can be exhibited without computing the group: $x \mapsto a(x - t)$ with a a quadratic residue sends (t, s) to $(0, a(s - t))$, and $a(s - t)$ sweeps the coset $\text{QR} \cdot (s - t)$, so one representative per coset suffices, the two cosets being mutually exclusive and exhaustive.

Corollary 2.2 (vertex anchoring). *Let $\text{Aut}(T)$ act on the vertices and choose one representative $v_1, \dots, v_r$ from each orbit. If T is k -inducible, then T has a witness in which some voter ranks some v_i first.*

Proof. The proof of Lemma 2.1 verbatim, with the action on vertices in place of the one on ordered pairs. ■

We say that a witness meeting either restriction is *anchored*. The two are alternatives rather than a sequence, each without loss of generality, and which is better depends on the tournament's number of vertex pair orbits vs vertex orbits, as well as the relative complexity of testing each restriction.

Lemma 2.1 must be given the full set of orbits. It licenses the restricted search *provided* the representatives meet every orbit. For every tournament we compute $\text{Aut}(T)$ separately, with **nauty** 2.8.6 [16], and record its orbits alongside the run. Every automorphism group, orbit set and canonical form reported in this paper comes from that same tool. Note that using the orbits of a known subgroup $\Gamma \leq \text{Aut}(T)$ is safe, since Γ -orbits are a refinement of the $\text{Aut}(T)$ -orbits with respect to the partition order.

Lemma 2.3 (voter lex-ordering). *If T is k -inducible, then T has a witness whose k orders are non-decreasing under any fixed total order on linear orders.*

Proof. The support $c(e)$ counts voters and is therefore invariant under permuting them, so any profile with the same multiset of orders induces the same tournament. Sorting the multiset gives the required witness. ■

The two lemmas are allowed to be applied simultaneously because the condition in Lemma 2.1 (and Corollary 2.2) is phrased as “*some* voter ranks ...”, not “voter 1 ranks ...”, making it invariant under permuting the voters. Both lemmas are therefore available to every formulation in Appendix A.

Together the two lemmas reduce the running time substantially. Lemma 2.3 is absorbed into the definition of a base state: the k orderings of B are sorted, thus dividing the count by a factor of almost $k! = 120$. Lemma 2.1 *invalidates* base states outright when *all* the representative pairs lie inside B : a base state can then extend to an anchored witness only if one of its voters already ranks a chosen pair at the top of B , a property of the base state alone. If even one representative lies outside B this test is unavailable, since a base state failing it may still extend to a witness anchored on that representative. We call a base state *live* if it respects Lemma 2.1. For instance, on the five-vertex base we use on the Paley tournaments this leaves 2,591 of the 8,031 base states live.

The search and the certification differ in *when* the condition of Lemma 2.1 is tested. Both test it at the start, which is what determines the *live* base states, those the anchoring does not discard. Our algorithmic approach then keeps testing it during the search, because the condition constrains the *full* orders while a base state fixes only their restriction to B : inserting a vertex above an anchored pair breaks it. Pruning on that is sound because a partial profile violating the condition has no extension satisfying it. The SAT certification instead tests it once, at the root, where it decides which cubes reach the solver.

2.3 The search in five steps

The search extends base states. A base state fixes how the k voters rank B ; the search inserts the remaining $n - |B|$ vertices into those rankings one at a time, backtracking as soon as the arcs already decided contradict T . A base state extended to k full rankings yields a witness; one that cannot be is eliminated; and T fails to be k -inducible exactly when every base state is eliminated. Algorithmically, with T and the margin regime as inputs, our approach has 5 steps:

1. **Pick a size.** Choose $b = |B|$, in practice 5 or 6.
2. **Pick a subtournament.** Choose B with $|B| = b$. Only $T|_B$ matters for what follows, so the choice can and should be optimised quickly (see Appendix A.2), since the base-state count varies by more than two orders of magnitude across the 12 isomorphism classes at $b = 5$.
3. **Find its base states.** Enumerate the k -tuples of linear orders of B whose supports agree with $T|_B$, counted up to voter permutation by Lemma 2.3.
4. **Determine the usable symmetries.** Compute $\text{Aut}(T)$, or a subgroup, and its orbits, then impose the ordered-pair anchoring of Lemma 2.1, the vertex anchoring of Corollary 2.2, or neither, discarding the base states no anchored witness can use.
5. **Run to completion, in parallel across base states.** The verdict is positive as soon as any base state yields a witness. It is negative only when *every* live base state has been searched to exhaustion.

Figure 1, in Appendix A.2, works all five steps through on P_7 at $k = 3$, where the witness is unique up to symmetry. Steps 1–4 are fast and user-chosen; step 5 is the parallel computation. The same five steps describe the SAT route, with step 5 replaced by handing each live base state to a CDCL solver as a *cube* (the partial assignment fixing the voters’ rankings of B), so the two share their decomposition and their two lemmas and differ only in what refutes a leaf in the search tree.

3. Results

The two regimes we explore are *unrestricted* (or plain majority) inducibility, and *unit margin*. The former is the standard notion of inducibility, and the latter asks that every arc be carried by exactly three voters against two. Only an unrestricted refutation bounds $N(5)$. However, the lemma below shows that for the Paley family and most tournaments derived from it, the unrestricted inducibility is equivalent to inducibility with margin at most 3 (i.e. support at most 4).

Lemma (3-cycle bound). *Let $u \rightarrow v \rightarrow w \rightarrow u$ be a directed triangle of T and let $\pi_1, \dots, \pi_k$ be any linear orders with k odd. Then*

$$c(u \rightarrow v) + c(v \rightarrow w) + c(w \rightarrow u) \leq 2k.$$

Consequently, in any majority witness, no arc lying in a directed triangle has unanimous support k .

Proof. A linear order cannot rank u above v , v above w and w above u simultaneously, so each voter agrees with at most two of the three arcs; summing over the k voters gives the inequality. If $c(u \rightarrow v) = k$, the other two supports sum to at most k , contradicting the majority requirement that each be at least $(k+1)/2$. $\square$

For $k = 3$ the second conclusion is Proposition 1 of Milosz, Hamel and Pierrot [27], with the same argument; we need it for every odd k , and in the summed form above.

3.1 Paley(23) and Paley(31) are not 5-inducible

Theorem. *P_{23} is not 5-inducible. Consequently $N(5) \leq 23$.*

The proof is a search that runs to completion. The base is $\{1, 2, 3, 6, 12\}$, which induces the regular tournament on 5 vertices and gives 8,031 base states. They are reduced to 2,591 by Lemma 2.1, by requiring one voter to rank either (1, 3) or (1, 6) as the first pair. None of them produces a witness. This statement is also formally certified in Appendix B.1.

The same method settles a larger case: P_{31} is not 5-inducible, over 21,009 base states with a different base, reduced to 4,007 by Lemma 2.1. Appendix A.3 shows how the runtime scales in q , and Appendix A.4 tabulates the cost of every computation in this paper. The value of settling P_{31} is to show that the bespoke approach decides 31-vertex instances on a single laptop within a day, something that SAT and ILP solvers are unable to do.

3.2 Criticality across the Paley family

Arc-criticality implies vertex-criticality: reversing an arc incident to v leaves $T - v$ unchanged, and the restriction of a witness to $V \setminus \{v\}$ is a witness for $T - v$. Each cell below therefore records only the *strongest* property established, at each of the two margins. “Not vertex-critical” therefore also rules out arc-criticality.

q	unit margin	unrestricted
19	arc-critical	5-inducible, so criticality does not apply
23	not vertex-critical	arc-critical (Section 3.5)
31	not vertex-critical	not vertex-critical (Section 3.4)
43	not vertex-critical	not vertex-critical (Section 3.4)

Note that P_{19} is the only member critical at unit margin, and it is critical in the strongest sense: reversing any one arc makes it unit-margin inducible, which by the implication also gives vertex-criticality. At the other end, P_{31} and P_{43} fail to be vertex-critical at *both* margins. The interesting cell is $q = 23$: it is *vertex-critical unrestricted*, but *not* vertex-critical at unit margin. In other words, $P_{23} - v$ is 5-inducible, but not with unit margin, analogous to the following result.

3.3 Paley(19) is 5-inducible, but not with unit margin

That P_{19} is 5-inducible is due to Bachmeier et al. [1], who computed $\dim(Q_{19}) = 5$. Our search recovers it independently and contributes the negative half: it is *not* 5-inducible with unit margin. This is a complete refutation over all 2,200 base states of the base $\{1, 2, 3, 4, 6\}$, and it is the second statement we certify formally in Appendix B.1. These two facts supply the separation in Section 1: no arc of P_{19} can be unanimous, by the

3-cycle bound above, so it is 5-inducible with margin at most 3 and not with unit margin, so the inclusion of the unit-margin regime in margin at most 3 is strict.

Both of our certified refutations, P_{23} at unrestricted margin in Section 3.1 and P_{19} at unit margin, are established formally by cube-and-conquer computations in which every leaf is refuted by CaDiCaL 2.0.0 [15, 20] in LRAT and rechecked by `lrat-trim` [20], a separate program that shares no code with the solver, with the solver’s own proof checking deliberately disabled. Both are complete on *both* halves (the leaves and the claim that the leaves are exhaustive), and both reduce to the same two human-checked lemmas, 2.1 and 2.3 (see Appendix B).

3.4 Paley(31) and Paley(43) are not vertex-critical

Vertex-criticality is the natural minimality notion for an *obstruction*, a tournament that is not 5-inducible, and P_{43} , the first explicit non-5-inducible tournament [2], is the natural candidate.

Theorem. *Neither P_{31} nor P_{43} is vertex-critical for 5-inducibility.*

The orbit break. For any $q \equiv 3 \pmod{4}$, $\text{Aut}(P_q - v)$ contains the stabiliser of v in $\text{Aut}(P_q)$, which is cyclic of order $(q - 1)/2$ and splits the remaining $q - 1$ vertices into exactly two orbits of size $(q - 1)/2$, the quadratic residues and the non-residues. Two representatives, one from each, therefore meet every orbit, and the vertex anchoring of Corollary 2.2 applies with those two.

Each sweep runs over all 8,031 base states of a five-vertex base inducing the regular tournament on five vertices: $\{1, 2, 3, 4, 13\}$ in $P_{31} - v$ and $\{1, 2, 3, 4, 11\}$ in $P_{43} - v$. The two counts coincide because, as in Section 2.1, the number of base states depends only on the induced subtournament and on the margin, and not on n .

Every arc of $P_q - v$ lies in at least $(q - 3)/4$ directed triangles, so no arc of a majority witness could be unanimous, and margin ≤ 3 is not a restriction at all. The refutation at margin ≤ 3 is therefore an unrestricted refutation. We use the same approach for both $q = 31$ and $q = 43$.

We thus independently reconfirm the main result of [2], that P_{43} is not 5-inducible: if it were, so would be $P_{43} - v$, by deleting v from every voter’s order.

3.5 Vertex and arc criticality for Paley(23)

Theorem. *P_{23} is arc-critical: reversing any single arc yields a 5-inducible tournament. Consequently it is vertex-critical.*

Let T^e denote P_{23} with the arc $(1, 2)$ reversed; since $\text{Aut}(P_{23})$ is transitive on arcs, any two single-arc reversals of P_{23} are isomorphic, so we may pick $(1, 2)$ without loss of generality. The computation is a scan of the 8,031 base states of T^e , but now T^e is rigid, so Lemma 2.1 and Corollary 2.2 do not apply. The witness we found was verified independently of the search that produced it, by two checks on the five orders: each is a permutation of the 23 vertices, and the majority tournament they induce agrees with T^e in every one of the 253 arcs and therefore differs from P_{23} in exactly e .

Vertex- and arc-criticality come apart, as far as they can. Arc-criticality implies vertex-criticality, so it is natural to ask whether the two coincide. They do not. Among the 110 vertex-transitive tournaments on 21 vertices, four are unit-margin obstructions, and all four are vertex-critical. However, arc-criticality takes all three of its possible values among those same four: two of them have every arc critical, one has exactly half, and one has none at all, so no statement of the form “this tournament is or is not arc-critical” captures the situation (full details in Appendix F).

3.6 Other structured tournament families

Beyond the Paley family we swept the structured families catalogued by McKay [26] exhaustively for $n < 23$. Every verdict is negative for the bound, so we only report them briefly.

First, *the vertex-transitive family contains no majority obstruction at all for $n < 23$* : all 110 are 5-inducible, 106 of them already at unit margin and the remaining four by exhibited majority witnesses. Since only a majority refutation can move $N(5)$, that family is closed as a source of candidates below 23, and the closure survives one-arc perturbation, since reversing any single arc of the four unit-margin obstructions, one representative per arc orbit, again leaves a majority-inducible tournament in all 40 cases.

Second, *both doubly regular tournaments on 19 vertices, P_{19} and the one other of that order, are unit-margin obstructions and both are majority-inducible*, so the separation that makes P_{19} interesting is not a property of P_{19} in particular. Nor is its criticality: the second one is *arc-critical* at unit margin too. Its automorphism group has order 3, so its 171 arcs fall into 57 orbits of size 3, and reversing the representative of any one of them yields a tournament with a unit-margin witness: 57 orbits and 57 witnesses. Arc-criticality at unit margin is therefore shared by both doubly regular tournaments of order 19.

Third, every regular tournament on at most 15 vertices is inducible at unit margin, over all $R_{15} \approx 1.8 \times 10^{10}$ of those on 15 vertices, and with the instance counts summing to OEIS A096368(7) [17]. Likewise, every self-converse tournament on 13 vertices is inducible at unit margin, over all $S_{13} \approx 9.5 \times 10^7$ of them. With the order-11 census of [2] and the order-12 analysis of Appendix C, that settles every self-converse tournament on at most 13 vertices.

Taken together with the Paley results above, these sweeps turn up ten tournaments that are not 5-inducible with unit margin: 2 doubly regular ones at order 19, 4 vertex-transitive ones at order 21, and 4 circulant ones at order 23.

One family needs no search. A tournament is *locally transitive* if every in- and out-neighbourhood induces a transitive subtournament. Every such tournament is k -inducible at unit margin for every odd $k \geq 3$, by an explicit construction rather than a computation; Appendix E provides it. This is the one structural family for which the inducibility question is settled outright.

4. Discussion and open problems

4.1 Other applications of incremental depth-first search with MRV

Subject to the caveats of Appendix A.4, the shape of the engine may transfer: a placement search in which each decision *refines*, rather than recomputes, the candidate sets of all remaining objects, with a dynamic minimum-remaining-values (MRV) order and an externally discharged symmetry break. The conditions are specific: the extensions of a partial solution must factor as a product that shrinks monotonically under each decision, and the symmetry group must be computable; a general-purpose solver may win if they do not apply.

Neither device is new in isolation, and the closest antecedent is Knuth’s Algorithm X with dancing links [18]. Its column-choice rule (take the column with the fewest remaining ones) is our MRV order, and its link surgery does what our refinement does: the state is altered in place and put back on backtrack, so a node costs what changed rather than what the whole state contains. The constraint is what differs. Exact cover partitions, whereas an arc of a tournament imposes a threshold, at least three of five voters, so the search here is not an exact cover instance and dancing links does not apply to it as it stands.

Two design lessons do seem to generalise. First, the incremental refinement, not the heuristic order, carried most of the speedup here, a factor of 110 against a factor of a few. Second, a symmetry break should be *parameterised* by its representatives rather than hard-coded, so that its correctness obligation becomes a separate computation instead of an assumption buried in the search; Section 3.4 is a case where discharging that obligation required an explicit orbit calculation that the engine could not have performed for itself.

4.2 Conjectures and open problems

Conjecture 1: The true value of $N(5)$ is 23, so every tournament of order less than 23 is 5-inducible.

Both ends of $13 \leq N(5) \leq 23$ are constructive and improved from [2]; the margin-1 obstructions at $n = 19$ and 21 cannot move the upper one, since every one of them is majority-inducible.

Conjecture 2: P_q is not 5-inducible for any eligible prime $q \geq 23$, i.e. $q \equiv 3 \pmod{4}$.

It holds at $q = 23, 31$ and 43 , and only finitely many q are in doubt: $\alpha^*(P_q) = \text{MAS}(P_q)/C$ tends to $\frac{1}{2}$, so the predictability eventually falls below $\frac{3}{5}$, past which the necessary condition of Section 1.2 refutes 5-inducibility with no search at all. What is missing is a bound on where that range ends; the values in Section 1.2 fall too irregularly to extrapolate one.

Open problem 1: Vertex- and arc-criticality. Whether P_q minus *two* vertices is 5-inducible, equivalently whether P_q minus one vertex is vertex-critical, is open at $q = 31$ and $q = 43$.

Open problem 2: The margin hierarchy. The first of unit margin $\subseteq$ margin $\leq 3 \subseteq$ unrestricted is strict (Sections 3.3 and 3.6); whether the second ever is, that is, whether some tournament is inducible by five voters but not at margin ≤ 3 , is open. The 3-cycle lemma forces such a tournament to have an arc in *no* directed triangle, which rules out every regular tournament. The conjecture of [2] is that the inclusion is strict at every odd $k > 3$.

5. Reproducibility

5.1 The repository

The code, data and verdict files are at <https://github.com/Leonardini/TournamentsBeyond5Voters>. Its `CLAIMS.md` carries one row per statement made here, naming the artifact that holds the verdict and the command that re-derives it, and it ends with an explicit list of what the package does not establish. A permanent archive of the commit corresponding to the accepted version will be deposited on acceptance and its DOI recorded here.

The repository contains the search engine as a single C translation unit, the SAT encoding and cube tooling in Python, the certification driver, the independent witness verifiers, the tournaments as bit strings, and a verdict record for every result reported here. Each verdict records the exact command line, the base and break used, the coverage audit, the runtime, and the two root hashes for the formally certified results.

5.2 What must and what may not replicate

The certified computations publish two root hashes, and only one of them is portable.

ROOT (CNF) commits to the SHA-256 of every cube’s CNF, in cube order. These files are generated by the repository’s own code from the tournament and the base, and are reproducible on any machine: a third party who regenerates the cube set *must* obtain this exact value. It is the meaningful cross-check, because it proves they solved the same problems.

ROOT (proofs) additionally commits to the LRAT proof bytes. These are *not* portable, as they depend on the exact SAT solver and version used. A mismatch here is not evidence of an error.

We have run that check: on the cluster, under a different compiler and CaDiCaL build, both refutations gave matching roots at 70.9 and 567.6 cluster core-hours against the 25.1 and 236.0 of Appendix A.4, and both coverage instances rebuilt to their committed CNF hashes and were refuted again.

Search runtimes: nodes replicate, hours do not. A three-way version of that distinction applies to the uncertified sweeps. *Verdicts* replicate under any correct implementation, which is the claim the paper makes. *Node counts* replicate exactly, but only with five choices held fixed:

1. the same tie-breaking rule in the variable order and in the domain refinement; ties are common and a different rule explores a different tree of the same size class but not the same size;
2. the same vertex numbering of the tournament;
3. the same base size and base;
4. the same resulting set of base states, which follows from (2) and (3);
5. the same margin regime.

The engine prints all of these except (2) in a header on every run (margin, variable order, tie-break mode, symmetry break, base size, base, base arc pattern and base-state count), so a run log is a self-describing fingerprint. Condition (2) is the tournament file itself, which is why the repository stores tournaments rather than regenerating them from their constructions. *Core-hours* replicate only in order of magnitude; Appendix A.3 measures a factor of 2 to 5 per core between the two machines we used.

Motivation and statement of AI use

This paper grew out of our previous paper [2], whose constructive bound $N(5) \leq 43$ came from a structural argument, and whose counting bound $N(5) \leq 38$ names no explicit tournament. The question of what a genuine decision procedure would reach, at the orders the screen cannot speak about, led to the engine of Appendix A.

During the preparation of this work, the authors used Claude for exploratory reasoning, implementation assistance, drafting and editing. The authors reviewed and edited the output and take full responsibility for the contents of the published article.

Acknowledgments

This work was granted access to the HPC resources of IDRIS under an allocation made by GENCI. LC acknowledges funding from the MRC Centre for Global Infectious Disease Analysis (reference MR/X020258/1), funded by the UK Medical Research Council (MRC). This UK funded award is carried out in the frame of the Global Health EDCTP3 Joint Undertaking.

References

- [1] G. Bachmeier, F. Brandt, C. Geist, P. Harrenstein, K. Kardel, D. Peters, H. G. Seedig. k -Majority digraphs and the hardness of voting with a constant number of voters. *Journal of Computer and System Sciences*, 105:130–157, 2019.
- [2] L. Chindelevitch, A. Harutyunyan. Tournaments determined by three and five voters. *Submitted*.
- [3] M. J. H. Heule, O. Kullmann, V. W. Marek. Solving and verifying the Boolean Pythagorean triples problem via cube-and-conquer. In *Theory and Applications of Satisfiability Testing (SAT 2016)*, LNCS 9710, pages 228–245, 2016. arXiv:1605.00723.
- [4] N. Alon, G. Brightwell, H. A. Kierstead, A. V. Kostochka, P. Winkler. Dominating sets in k -majority tournaments. *Journal of Combinatorial Theory, Series B*, 96(3):374–387, 2006.

- [5] R. L. Graham, J. H. Spencer. A constructive solution to a tournament problem. *Canadian Mathematical Bulletin*, 14(1):45–48, 1971.
- [6] C. Eggermont, C. Hurkens, G. J. Woeginger. Realizing small tournaments through few permutations. *Acta Cybernetica*, 21(2):267–271, 2013.
- [7] D. Shephardson, C. A. Tovey. Smallest tournaments not realizable by $\frac{2}{3}$ -majority voting. *Social Choice and Welfare*, 33(3):495–503, 2009. doi:10.1007/s00355-009-0375-7.
- [8] B. Dushnik, E. W. Miller. Partially ordered sets. *American Journal of Mathematics*, 63(3):600–610, 1941.
- [9] M. J. H. Heule, O. Kullmann, S. Wieringa, A. Biere. Cube and conquer: guiding CDCL SAT solvers by lookaheads. In *Haifa Verification Conference (HVC 2011)*, LNCS 7261, pages 50–65, 2012.
- [10] L. Cruz-Filipe, M. J. H. Heule, W. A. Hunt Jr., M. Kaufmann, P. Schneider-Kamp. Efficient certified RAT verification. In *Automated Deduction (CADE 2017)*, LNCS 10395, pages 220–236, 2017. This is the paper introducing the LRAT format that `lrat-trim` consumes.
- [11] L. Cruz-Filipe, J. Marques-Silva, P. Schneider-Kamp. Formally verifying the solution to the Boolean Pythagorean triples problem. *Journal of Automated Reasoning*, 63(3):695–722, 2019. doi:10.1007/s10817-018-9490-4.
- [12] M. J. H. Heule. Everything’s bigger in Texas: the largest math proof ever. ACL2 Seminar, University of Texas at Austin, 9 September 2016. Joint work with O. Kullmann and V. W. Marek. <https://www.cs.utexas.edu/~moore/acl2/seminar/2016.09.09-heule/ACL2.pdf> Cited for the compressed size of the published cube certificate; the re-verification runtimes quoted in Appendix B.1 are stated in [3] itself.
- [13] M. J. H. Heule. Schur number five. In *AAAI 2018*, pages 6598–6606, 2018.
- [14] J. Brakensiek, M. Heule, J. Mackey, D. Narváez. The resolution of Keller’s conjecture. In *Automated Reasoning (IJCAR 2020)*, LNCS 12166, pages 48–65, 2020.
- [15] A. Biere, K. Fazekas, M. Fleury, M. Heisinger. CaDiCaL, Kissat, Paracooba, Plingeling and Treengeling entering the SAT Competition 2020. In *Proceedings of SAT Competition 2020: Solver and Benchmark Descriptions*, volume B-2020-1 of *Department of Computer Science Series of Publications B*, pages 50–53. University of Helsinki, 2020. The solver description; for the LRAT support we use, and for `lrat-trim`, see [20].
- [16] B. D. McKay, A. Piperno. Practical graph isomorphism, II. *Journal of Symbolic Computation*, 60:94–112, 2014. The `nauty` and `gentourng` tools, used for all tournament enumeration.
- [17] OEIS Foundation Inc. Entry A096368, *Number of unlabeled regular tournaments with $2n + 1$ nodes*. The On-Line Encyclopedia of Integer Sequences, <https://oeis.org/A096368>.
- [18] D. E. Knuth. Dancing links. In J. Davies, B. Roscoe, J. Woodcock, editors, *Millennial Perspectives in Computer Science*, pages 187–214. Palgrave, 2000. arXiv:cs/0011047.
- [19] N. Wetzler, M. J. H. Heule, W. A. Hunt Jr. DRAT-trim: efficient checking and trimming using expressive clausal proofs. In *Theory and Applications of Satisfiability Testing (SAT 2014)*, LNCS 8561, pages 422–429, 2014. doi:10.1007/978-3-319-09284-3_31.
- [20] F. Pollitt, M. Fleury, A. Biere. Faster LRAT checking than solving with CaDiCaL. In *Theory and Applications of Satisfiability Testing (SAT 2023)*, LIPIcs 271, pages 21:1–21:12, 2023. doi:10.4230/LIPIcs.SAT.2023.21. Introduces `lrat-trim` and CaDiCaL’s native LRAT output, which is the capability we use; it postdates the solver description [15].
- [21] OEIS Foundation Inc. Entry A000568, *Number of outcomes of unlabeled n -team round-robin tournaments*. The On-Line Encyclopedia of Integer Sequences, <https://oeis.org/A000568>.
- [22] OEIS Foundation Inc. Entry A002785, *Number of self-complementary oriented graphs with n nodes*. The On-Line Encyclopedia of Integer Sequences, <https://oeis.org/A002785>. McKay’s note on that entry records that these are also the self-converse tournaments.

- [23] A. E. Brouwer. The enumeration of locally transitive tournaments. Report ZW 138/80, Mathematisch Centrum, Afdeling Zuivere Wiskunde, Amsterdam, 1980.
- [24] B. Alspach, C. Tabib. A note on the number of 4-circuits in a tournament. *Annals of Discrete Mathematics*, 12:13–19, 1982.
- [25] B. Debord. Caractérisation des matrices des préférences nettes et méthodes d’agrégation associées. *Mathématiques et sciences humaines*, 97:5–17, 1987.
- [26] B. D. McKay. Combinatorial data: digraphs. <https://users.cecs.anu.edu.au/~bdm/data/digraphs.html>. Source of the circulant, vertex-transitive and doubly regular tournament catalogues swept in Section 3.6.
- [27] R. Milosz, S. Hamel, A. Pierrot. Median of 3 permutations, 3-cycles and 3-hitting set problem. In *Combinatorial Algorithms (IWOCA 2018)*, LNCS 10979, pages 224–236, 2018. doi:10.1007/978-3-319-94667-2_19.

Appendix A. The search engine

Both integer programming and satisfiability are central to this project rather than foils for it. Every computational claim in the previous paper [2] is certified by an integer or linear program, and the integer program of A.1 below is that paper’s formulation B.1, verbatim; the only machine-checked theorem in this paper is a SAT computation from end to end, with CaDiCaL [15] emitting LRAT proofs that `lrat-trim` checks. The claim here is narrow: on this one decision problem, at this one range of sizes, neither tool reaches the answer while a search built around the problem’s own structure does, and the same SAT machinery is then indispensable for turning that answer into a proof.

A.1 Integer programming and SAT are both too slow

The natural integer program has one binary $x_{v,ij}$ per voter v and pair $i < j$, meaning that v ranks i above j , with $x_{v,ji}$ read as $1 - x_{v,ij}$ so that antisymmetry is structural. Each voter is forced to be a linear order by forbidding both orientations of every 3-cycle,

$$0 \leq x_{v,ij} + x_{v,jl} - x_{v,il} \leq 1 \quad (i < j < l),$$

and the profile is forced to induce T by one majority row per arc, $\sum_v x_{v,ij} \geq (k+1)/2$. This is a pure feasibility problem. At $q = 23$, $k = 5$ it has 1,265 binaries, 17,710 transitivity rows and 253 majority rows.

The obstruction is structural rather than a deficiency of the solver. Put $x_{v,ij} = 3/5$ on every arc of T and $2/5$ on every non-arc, the same for all voters. Every cyclic triple then reads $3/5 + 3/5 - 2/5 = 4/5$ and every transitive one $3/5$, both interior to $[0, 1]$, while every majority row is tight at exactly 3. The relaxation is therefore feasible, and since the problem carries no objective, LP feasibility yields no bound whatever: branch-and-bound is left with nothing to prune by except infeasibility discovered deep in the tree. That fractional point is also invariant under both symmetry groups acting on the problem (the $k!$ voter relabellings and $\text{Aut}(T)$), so it supplies no branching direction either, and orbital branching recovers only the $k!$ factor. CPLEX and Gurobi both failed to close instances at $n = 19$ within a day.

The satisfiability encoding is *dissent-Boolean*: one variable per voter v and arc $i \rightarrow j$ of T , true when v ranks j above i and so dissents from that arc. Each arc bounds its dissents, at most $(k-1)/2$ of them for majority and exactly $(k-1)/2$ for unit margin, and each voter’s dissent set is forced to leave a linear order triple by triple: a tournament is transitive exactly when it has no directed 3-cycle, so it is enough to forbid, for every voter and every triple, the two dissent patterns that would make that triple cyclic. For $k = 5$ that is $5\binom{n}{2}$ variables, 15 clauses per arc at unit margin and 10 per triple. It is compact and propagates well locally, and it is what both certified refutations of Appendix B are built on.

Table 4 of [1] reports their **Sat-Check-k-Majority** on uniform random tournaments, tabulating only entries whose average stays under 20 seconds. The $k = 5$ column reads

n	18	19	20	21	22	23	24
seconds	0.23	0.35	0.54	5.87	11.07	18.95	—

so on *random* instances the method leaves the table between $n = 23$ and $n = 24$, roughly doubling per added vertex. Random tournaments are the easy case: no automorphisms to rediscover and, being far from extremal, many witnesses. On the structured instance of the same order, P_{23} , the same solver did not terminate within a cumulative six weeks. The gap between 18.95 seconds and six weeks at equal n is the cost of symmetry.

A better encoding and an explicit symmetry break both help substantially and are still not enough (results not shown). What does work is supplying the solver with the decomposition itself: cube-and-conquer over base states, which is how both refutations are certified in Appendix B. A CDCL solver must otherwise discover that decomposition for itself, and cannot exploit $\text{Aut}(T)$ at all without an explicit break in the encoding.

A.2 Depth-first placement search

The search maintains a partial profile: a set $S \subseteq V$ of *placed* vertices together with, for each voter, a linear order of S . It begins from a base state, which places B , and extends by inserting one unplaced vertex at a time. Inserting v means choosing, for each voter independently, a slot among the $|S| + 1$ available; the *domain* $D(v \mid S)$ is the set of k -tuples of slots consistent with every arc between v and S . A partial profile is extendable only if every unplaced vertex has a non-empty domain, and domains shrink monotonically as S grows, which makes the search a constraint-propagation problem rather than a blind enumeration.

ALGORITHM 1. Placement search for k -inducibility of T

Input : tournament T on vertex set V ; number of voters k ; margin regime (majority or unit)
Output: a witness profile, or a refutation

MAIN

```

for each base state (a lex-canonical assignment of the base  $B$  to  $k$  orders):
     $S \leftarrow B$ 
     $D(v \mid S) \leftarrow$  fresh enumeration, for every  $v$  not in  $S$ 
    if any  $D(v \mid S)$  is empty: continue // this base state is dead
    if DFS( $S$ ) succeeds: return its witness
return "not  $k$ -inducible" // no base state extends

```


DFS(S)

```

if  $S = V$ : return the profile // all vertices placed
 $v^* \leftarrow \text{argmin } |D(v \mid S)|$  over  $v$  not in  $S$  // MRV, recomputed at EVERY node
    ties broken toward  $\max | \text{out}(v) \ \& \ S | - | \text{in}(v) \ \& \ S | |$ 
for each slot tuple  $p$  in  $D(v^* \mid S)$ :
    insert  $v^*$  at  $p$  in each of the  $k$  orders;  $S' \leftarrow S \cup \{v^*\}$ 
    for each  $w$  not in  $S'$ :
         $D(w \mid S') \leftarrow \text{REFINE}(D(w \mid S), p, w, v^*)$ 
        if  $D(w \mid S')$  is empty: undo; continue with the next  $p$ 
    if DFS( $S'$ ) succeeds: return its witness
    undo // arena pointer reset
return failure

```

REFINE(D , q , w , u) // q = the slot tuple just used for u

(a) the tournament P_7 

(b) the base state, forced

	0	1	2	3
π_1	1	2	4	
π_2	2	4	1	
π_3	4	1	2	

the only one

options, as slot triples:

$D(3)$	113	203	212	213	230
	$ D(3) = 5$	fewest, so $v^* = 3$			
$D(5)$	032	122	131	132	302
	$ D(5) = 5$				
$D(6)$	023	221	311	320	321
	$ D(6) = 5$				
$D(7)$	000	001	002	003	010
	012	020	030	100	101
	111	120	200	201	300
	$ D(7) = 17$				

(c) two rejected insertions of 3

put 3 at 113:

π_1	1	3	2	4
π_2	2	3	4	1
π_3	4	1	2	3

breaks Lemma 2.1: no voter can still lead with (1, 2) or (1, 4)

put 3 at 213:

π_1	1	2	3	4
π_2	2	3	4	1
π_3	4	1	2	3

$D(5 \mid B \cup \{3\}) = \emptyset$; and it makes $2 \rightarrow 3$ unanimous

(d) an insertion of 3 that survives

put 3 at 230:

π_1	1	2	3	4
π_2	2	4	1	3
π_3	3	4	1	2

π_1 still leads with (1, 2), and nothing has run out of room

what that leaves of (b)'s options:

$D(5)$	043	142	143	403
	$ D(5) = 4$	was 5		
$D(6)$	024	222	420	
	$ D(6) = 3$	was 5		
$D(7)$	041	301	302	401
	$ D(7) = 4$	was 17		

(e) refining $D(6)$

past 230:

tuple	per voter	lift	$c_6(3)$
023	=, =, +1	→ 024	2 ✓
221	split, =, +1	→ 222	2 ✓
		→ 322	1 ✗
311	+1, =, +1	→ 412	1 ✗
320	+1, =, split	→ 420	2 ✓
		→ 421	1 ✗
321	+1, =, +1	→ 422	1 ✗

(f) the completed profile

π	1	2	3	4	5	6	7
$g\pi$	5	7	2	4	6	1	3
$g^2\pi$	6	3	7	4	1	5	2

$g = (156)(273)$, of order 3
every arc at support 2

Figure 1: Algorithm 1 on P_7 at $k = 3$, which is 3-inducible, so the run ends in a witness. A *slot triple* is one gap per voter, of the $4^3 = 64$ available, and a vertex's domain is those triples its arcs to B admit. (a) The base $T|_B$ on $B = \{1, 2, 4\}$ in bold; dashed, the arcs at the first vertex inserted. (b) Steps 2–4: the base state is forced, and MRV takes the smallest domain. (c) Two insertions the search rejects, one by Lemma 2.1 and one by leaving 5 no slot at all. (d) One that survives, and what it leaves of (b). (e) REFINES lifts each parent tuple past the slot just used (unchanged below it, shifted above it, split where the two coincide, which is what green marks), and the one new arc then keeps or kills each lift. (f) The witness, every arc at support exactly 2. Without Lemma 2.1 the search returns all 7 witnesses, a single orbit of $\text{Aut}(P_7)$, of order 21, so the profile is unique up to symmetry, as Shepardson and Tovey [7] report; π generates the other two under the g shown, the map $x \mapsto 2x + 4$ of $\mathbb{F}_7$.

```

D' <- {}
for each p in D:
  lift p to p' voter by voter:
    p_i < q_i -> p'_i = p_i           // w stays before u
    p_i > q_i -> p'_i = p_i + 1       // w stays after u, shifted by the insertion
    p_i == q_i -> BOTH are legal, so the tuple splits in two
  keep those lifts whose support c_w(u) = #{ i : w before u }
  meets the requirement of the single new arc between w and u
return D'

```

Four choices make it fast, and the second is the one that made $n = 19$ and $n = 23$ decidable at all.

A dynamic variable order. At every node we insert the unplaced vertex minimising $|D(v \mid S)|$, breaking ties toward the vertex whose arcs to S are most imbalanced. The important part is that the order is recomputed at every node rather than fixed per base state: on Paley instances a static order is not competitive, because

how constrained a vertex is depends strongly on which part of the tournament has already been placed.

Incremental domain refinement. Recomputing $D(v \mid S)$ from scratch at each node is the dominant cost of a naive implementation; instead the child’s domain is derived from the parent’s by the three cases in REFINE above, after which the single new arc filters the lifted tuples by their support. Child domains are therefore refinements of parent domains, computed in time proportional to the surviving tuples rather than to the whole slot space, and domains live in a stack-disciplined arena indexed by (depth, vertex) so that backtracking is a pointer reset. Like for like this is worth a factor of about 110 in wall time.

Symmetry breaking. The $k!$ voter symmetry is absorbed exactly by the base-state definition, which counts non-decreasing tuples only. For $\text{Aut}(T)$ the engine imposes the anchoring of Lemma 2.1 or of Corollary 2.2, whichever is stronger on the tournament: the condition is invariant under permuting voters, so it composes with the base lex-order, and monotone under insertion, so it can be tested during the search rather than only at the leaves.

Early termination and caching by arc pattern. A quick per-voter agreement bound, obtained by counting the agreement between voter i and the arcs between v and S two ways, proves $D(v \mid S)$ empty in $O(k|S|)$ time without enumerating it. And the per-base-state work is keyed by the *arc pattern* of $T|_B$ rather than by B itself, so base-state counts, which depend only on that pattern, are reused across every copy of the same base class and across every q .

Choosing the base. The base set is not a free parameter, and the fewest live base states is not the fastest choice: at $q = 31$ the class with 4,007 live states beat the class with 2,537, whose states are individually slower, so we rank the twelve 5-vertex classes by the product of the live count and a sampled time per state rather than by the count alone.

A.3 How the runtime scales in q

A note on $q = 27$. The series below includes $q = 27$, the one prime power among the orders we computed with. Since $27 = 3^3$, P_{27} must be built over $\mathbb{F}_{27}$ and not over $\mathbb{Z}/27$. P_{27} is the unique most symmetric doubly regular tournament on 27 vertices, with $|\text{Aut}(P_{27})| = 1053 = 3q(q-1)/2$, because the field automorphisms contribute (in general $|\text{Aut}(P_q)| = q(q-1)e/2$ when $q = p^e$). It is still transitive on arcs and on non-arcs, which is what Lemma 2.1 needs. We ran $q = 27$ to completion and it is not 5-inducible; since that does not bear on $N(5)$ we report it here only as a runtime datapoint.

On a fixed base decomposition the total runtime *decreases* as the tournament grows, and so does the node count: a larger Paley tournament is *more* constrained, so the search is refuted nearer the root and the tree above each base state is smaller rather than larger. The work per node moves the other way, rising monotonically from 10.8 to 57.4 microseconds across the four runs below, because a larger tournament carries larger domains and each refinement step costs more. The node count falls faster than the per-node cost rises, which is what makes the total fall. The full runs:

q	live base states	core-hours	nodes	seconds per live base state
23	2,591	34.03	1.14×10^{10}	47.3
27	2,537	31.04	6.22×10^9	44.0
31	4,007	26.89	3.43×10^9	24.2
43	2,537	22.87	1.43×10^9	32.5

P_{43} is decided in *less time than* P_{27} , and explores 4.3 times fewer nodes to do it. The last column is the third divided by the second, so only the core-hours and the live base states are measured. Just two of these four runs share a decomposition, $q = 27$ and $q = 43$, on the same five-vertex base class and the same 2,537 live base states; $q = 23$ and $q = 31$ each used their own best configuration, and $q = 31$ ’s is the cheaper class of Appendix A.2, which is why its per-state figure is the lowest of the four. The core-hours fall monotonically in q either way. Note, however, that P_{43} minus a vertex takes 185.5 core-hours, far more than P_{43} itself, as vertex deletion shrinks the automorphism group.

Additionally, the search gets faster as the margin is tightened. On the identical tournament ($P_{43} - v$), engine and 5-vertex base class, there is a factor of about 90. Requiring every arc to be at unit margin admits 2,200 base states where majority admits 8,031; additionally, under exact margin the last voter is *forced*, since the support accumulated over the first $k - 1$ voters determines what the last must contribute, so a whole level of the search collapses.

Two machines, and why their runtimes are not comparable. Some computations ran on a cluster rather than the laptop. Measured across six independent experiments, the cluster ran at 2 to 5 times the laptop’s runtime per unit of work.

We are not claiming that bespoke search beats integer programming or satisfiability in general. Our engine is narrow: it exploits the specific facts that a partial profile’s extensions form a product of per-voter slot choices, that this product refines monotonically, and that $\text{Aut}(T)$ is large and computable. For the specific Paley instances in this paper that the general-purpose methods did not settle for us, the search settles in hours on a single laptop.

A.4 Runtimes for every computation

All runs in this manuscript were performed on a single laptop (Apple M-series, 14 cores, 24 GB RAM), using at most 11 cores concurrently, with three exceptions that used a compute cluster: the census of regular tournaments on 15 vertices, the self-converse census of order 13, and the order-12 census of Appendix C. Runtimes are reported in *core-hours*, the workers’ running times summed.

computation	margin	verdict	core-hours
P_{19} , certified refutation	unit	not inducible	25.1
P_{23} , certified refutation	unrestricted	not inducible	236.0
$P_{31} - v$, sweep	unrestricted	not inducible	239.5
P_{23} , one arc reversed	unrestricted	inducible	210.4
second doubly regular 19, all 57 arc reversals	unit	all inducible	3.59
regular tournaments on 15 vertices	unit	all inducible	3,106
self-converse tournaments on 13 vertices	unit	all inducible	1,310
all tournaments on 12 vertices	unit	all inducible	$\approx 10,000$

The four full Paley sweeps, $q = 23, 27, 31$ and 43 , are tabulated with their node counts in A.3 above.

Appendix B. The formal certificates

B.1 The certified refutations

Both machine-checked results follow an established design rather than a new one. *Cube-and-conquer* is due to Heule, Kullmann, Wieringa and Biere [9]: split a hard instance into many partial assignments (*cubes*), hand each to a CDCL solver, and let the splitting supply the global structure the solver cannot find for itself. The standard of *proof* is the one Heule, Kullmann and Marek set for the Boolean Pythagorean triples problem [3], and Schur number five [13] and the Keller conjecture [14] after it: every subproblem emits a machine-checkable refutation, validated by an independent checker rather than trusted. We claim no methodological novelty.

Two differences are consequences of our setting. Our cubes are *base states*, a combinatorial decomposition, which lets the coverage of the cube set be certified separately (B.2) instead of resting on a splitting heuristic. And where [3] published a cube certificate from which its refutations can be regenerated, we *verify and discard*: each LRAT proof is checked, hashed and deleted, and we only publish the hash value. The choice suits the scale. Verifying the Boolean Pythagorean triples proof costs about 13,000 core-hours to regenerate the refutations from the published cube certificate [12], and about 16,000 to check them [3]. Our certificates are

orders of magnitude smaller: 25.1 core-hours for P_{19} and 236.0 for P_{23} cover solving *and* checking. Portable root values, which is exactly the check Section 5.2 supports, can be easily compared by anyone replicating our computation.

For P_{19} at unit margin the instance is the dissent-Boolean encoding of Appendix A.1 at $q = 19$, $k = 5$: 855 variables and 12,255 clauses. The cubes are the 142,251 base states of the 6-vertex base $\{1, 2, 3, 4, 5, 12\}$, of which 22,876 survive the arc-orbit break, and each is handed whole to CaDiCaL with `--lrat=true --checkproof=0`. Each LRAT proof is checked by `lrat-trim`. The runtime was 2.64 hours on ten cores. Its coverage half is the same construction as the one below, at 143,032 clauses, and solves in 7.56 seconds.

For the P_{23} refutation of Section 3.1, the base is $\{1, 2, 3, 4, 6, 7\}$ with 3,414,729 base states, of which 343,896 survive the arc-orbit and non-arc-orbit breaks; every one is UNSAT with an LRAT proof independently verified by `lrat-trim`. The runtime was 21.72 hours on eleven cores. The coverage half of the proof is a single unsatisfiable instance of 350 constraint clauses, 356 symmetry-breaking clauses and 3,414,729 negated cubes, one per base state, since coverage must see *every* base state whereas the orbit breaks apply to the search half only. It solves in just under 20 minutes on one core and emits a 2.05 GiB LRAT proof.

B.2 What the certificates do and do not establish

The trust chain has four links:

link	status
every live cube is unsatisfiable	machine-checked: CaDiCaL, rechecked by <code>lrat-trim</code>
the cubes cover every lex-canonical base assignment	machine-checked: LRAT, verified by <code>lrat-trim</code> ; also as DRAT by <code>drat-trim</code> [19]
(L1) orbit anchoring loses no generality	human-checked: Lemma 2.1
(L2) lex-ordering the voters loses no generality	human-checked: Lemma 2.3

The two human-checked links are proved in Section 2. Lemma 2.1 is what licenses solving only the *live* cubes, in its arc/non-arc form: $\text{Aut}(P_q)$ is transitive on arcs and on non-arcs, so a voter’s top pair may be assumed to be a fixed arc or a fixed non-arc. Lemma 2.3 licenses the lex chain in the coverage computation. They are human-checked links not because they are doubtful but because nothing in the pipeline checks them; the depth-first engine inherits them, but shares no implementation with this one, so their agreement is evidence about the implementations, not the lemmas.

Appendix C. The order-12 case analysis

Every tournament on 12 vertices is 5-inducible, so $N(5) \geq 13$. This appendix describes that computation. It is one of the results in this paper that was not run on the laptop of Section 3: it ran on a shared cluster, and its runtime is reported in cluster core-hours (see Appendix A.3).

C.1 Reduction to one-vertex extensions

A direct census is out of reach: $D_{12} \approx 1.5 \times 10^{11}$ classes, 170 times the order-11 census of [2]. We use that census as the outer loop instead. Let us write $L + S$ for the tournament obtained from an order-11 tournament L by adding a vertex v whose out-neighbourhood is $S \subseteq V(L)$. Two observations simplify the analysis.

(C1) Every order-12 tournament is $L + S$ for some order-11 L and some S .

(C2) Restricting a profile that realises $L + S$ to $V(L)$ gives a profile that realises L .

By (C1), enumerating all D_{11} order-11 classes and settling all $2^{11} = 2,048$ extensions of each settles every order-12 tournament; all of the former are 5-inducible by the census of [2], which is what makes the outer loop available. By (C2), a profile realising $L + S$ may be sought by *inserting* v into a profile realising L with no loss of generality, since every profile for $L + S$ arises that way. That is what phase 1 exploits.

C.2 Three phases

The census is organised into three phases, with the fastest ones occurring earlier.

Phase 1, the covering search. Rather than 2,048 separate instances per order-11 tournament, one search enumerates profiles realising L at unit margin and, for each, records every out-neighbourhood obtainable by inserting v into its five orders: v 's position in each order is a threshold, and the support of $v \rightarrow u$ is the number of orders that place v above u , so one traversal certifies many extensions at once. This is the engine of Appendix A.2 with v 's placement read off rather than fixed in advance. The margin makes this phase pay twice: it keeps the enumeration of profiles for L small, and a unit-margin witness is in particular a majority witness, so every extension the phase covers is settled in both regimes at once. The phase is one-sided: an extension it covers is certified 5-inducible, but one left over is not refuted. This phase settles every one of the 2,048 extensions of 99.5% of the order-11 tournaments.

Phase 2, deciding what is left. Each extension phase 1 leaves open is materialised as an order-12 tournament $L + S$ and decided by the same engine, again at unit margin, where a witness settles it positively. About one extension in 200,000 reached this phase, and every one of them is 5-inducible with unit margin.

Phase 3, plain majority. Only a *majority* refutation would be a counterexample, so any instance negative at unit margin is re-run without the margin restriction. Phase 3 was never invoked: every tournament was settled at unit margin, so the weaker question never had to be asked.

C.3 One tournament of each converse pair

If $\pi_1, \dots, \pi_5$ realise T , their reversals realise the converse $\bar{T}$, with the same margins, so T is 5-inducible if and only if $\bar{T}$ is. Moreover $\overline{L+S} = \bar{L} + (V(L) \setminus S)$, so settling all 2,048 extensions of L settles all 2,048 of $\bar{L}$ as well. We therefore keep one tournament of each converse pair, by the rule that L is kept if and only if $\text{canon}(L) \leq \text{canon}(\bar{L})$. This keeps exactly one of each pair together with every self-converse tournament, $(D_{11} + S_{11})/2$ in all, where S_{11} is the number of self-converse order-11 classes; this saves a factor of $\lesssim 2$. We also check the number of tournaments kept is $(D_{11} + S_{11})/2$.

Total runtime: approximately 10,000 cluster core-hours. The search code, the per-part records and the roll-up are at `cluster/jz_n12cover` in the repository of Section 5.

Appendix D. Exact enumeration counts

The counts the main text gives to two significant figures, in full. D_n , R_n and S_n are the numbers of isomorphism classes of all, of regular, and of self-converse tournaments on n vertices. They match OEIS A000568 [21], A096368 [17], and A002785 [22].

n	D_n	R_n	S_n
11	903,753,248	1,223	279,968
12	154,108,311,168	—	1,492,288
13	—	1,495,297	95,458,560
15	—	18,400,989,629	—

Appendix E. Locally transitive tournaments, at unit margin, for every odd $k \geq 3$

A tournament is *locally transitive* if for every vertex v both $N^+(v)$ and $N^-(v)$ induce transitive subtournaments. Equivalently it is *round*: the vertices admit a cyclic order in which every $N^-(v) \cup \{v\}$ and $\{v\} \cup N^+(v)$ is an interval. The equivalence is due to Brouwer [23], who proved it in the course of enumerating the class, and to Alspach and Tabib [24], who state it for tournaments under the name *domination orientable*; Bang-Jensen later generalised it to digraphs.

Theorem E.1. *Every locally transitive tournament is 3-inducible with unit margin, and hence k -inducible with unit margin for every odd $k \geq 3$.*

Only the strongly connected case arises, because a locally transitive tournament that is not strong is transitive: a strong component C with $|C| \geq 3$ contains a directed triangle, and any vertex outside C either beats all of C or loses to all of it, which would put that triangle inside its out- or in-neighbourhood. A transitive tournament is 1-inducible, and this profile can be padded by any linear order together with its reverse to achieve margin 1 with $k = 3$ and, by induction, any other odd k .

Let T be strong and round, cut the cyclic order at v_0 to get $A = (v_0, \dots, v_{n-1})$, write d_i for the out-degree of v_i and $r_i = i + d_i$ for the position of its last out-neighbour before reduction. The round condition gives $j \leq d_i \iff d_{i+j} \leq n - 1 - j$ for $1 \leq j \leq n - 1$, from which $r_0 \leq r_1 \leq \dots \leq r_{n-1}$: the arcs A ranks *wrongly* form a nested staircase. Its sources $H = \{i : r_i \geq n\}$ are a suffix $\{q, \dots, n - 1\}$, the ceilings $t_i = r_i - n$ satisfy $t_q \leq \dots \leq t_{n-1}$, and v_i 's back-targets are exactly $v_0, \dots, v_{t_i}$. Writing $L = \{0, \dots, q - 1\}$, one checks $t_{n-1} < q$, so H and the targets are disjoint, and for $h \in H, l \in L$ the round rule reads $v_h \rightarrow v_l \iff l \leq t_h$. Take

$$A = (v_0, \dots, v_{n-1}), \quad B = (v_q, \dots, v_{n-1}, v_0, \dots, v_{q-1}),$$

and let C list the vertices by *decreasing* key, where $\text{key}(v_l) = l$ for $l \in L$ and $\text{key}(v_h) = t_h + \frac{1}{2}$ for $h \in H$, ties among equal t_h broken by larger index first. The integer and half-integer keys never collide, so only that one tie-break is at issue, and by construction C ranks v_h before v_l exactly when $l \leq t_h$, that is exactly when $v_h \rightarrow v_l$.

Every arc is then ranked correctly by exactly two of the three. An arc inside L or inside H is ranked correctly by A and by B , and wrongly by C : inside L because the keys are the indices and C reverses them, inside H because the ceilings are monotone, so the head's key is at least the tail's and the tie-break settles equality the same way. A forward arc from L to H is correct in A and C and wrong in B ; a back arc from H to L is correct in B and C and wrong in A . Every arc therefore has support exactly 2, none is unanimous, and the profile realises T at unit margin; the padding argument above extends this margin to any odd $k \geq 3$. That *some* profile realises a tournament with all margins equal to 1 is a consequence of Debord's parity theorem [25]; the content of Theorem E.1 is that three voters suffice, uniformly over the family.

Appendix F. Criticality at 21 vertices

This appendix carries the detail behind Section 3.5. Throughout, h ranges over the four unit-margin obstructions among the 110 vertex-transitive tournaments on 21 vertices (Section 3.6).

Vertex-critical but maximally not arc-critical. Arc-criticality implies vertex-criticality, and the converse fails here as widely as it can. Each h is *vertex-critical* at unit margin, and for h_1 every single-arc reversal remains a unit-margin obstruction.

Arc-criticality is not a property of the tournament. All four have ten arc orbits of size 21; we sweep one representative per orbit at unit margin and get results that differ by tournament *and* by arc, U marking a reversal that is still not unit-margin inducible and S one that is:

tournament	Aut	the ten orbits	obstructions
h_1	$\mathbb{Z}_{21}$	U U U U U U U U U U	10
h_2	$\mathbb{Z}_{21}$	U S U S S U U S S U	5
h_4, h_5	$\mathbb{Z}_7 \rtimes \mathbb{Z}_3$	S S S S S S S S S S	0

Two things follow. First, arc-criticality at unit margin is *determined by the arc* and not by the tournament: h_2 splits five orbits to five, so exactly half its arcs are critical and we call it *arc-semi-critical*; with P_{23} arc-critical unrestricted and h_1 critical nowhere, all three possibilities occur among our tournaments. Second, every one of the 15 reversals that remains an obstruction belongs to a tournament with *cyclic automorphism group*, while all 20 orbits of the two with the nonabelian group of order 21 are inducible.